



**Министерство науки и высшего образования Российской Федерации
Федеральное государственное бюджетное образовательное учреждение
высшего образования
«Московский государственный технический университет
имени Н.Э. Баумана
(национальный исследовательский университет)»
(МГТУ им. Н.Э. Баумана)**

**Факультет «Информатика и системы управления»
Кафедра ИУ5 «Системы обработки информации и управления»**

**Отчет
по дисциплине «Технология Машинного обучения»**

**Выполнил:
студент группы ИУ5-62
Миронов Святослав
подпись, дата**

2019 г.

Задание:

1. Выберите набор данных (датасет) для решения задачи классификации или регрессии.
2. В случае необходимости проведите удаление или заполнение пропусков и кодирование категориальных признаков.
3. С использованием метода `train_test_split` разделите выборку на обучающую и тестовую.
4. Обучите две ансамблевые модели. Оцените качество моделей с помощью одной из подходящих для задачи метрик. Сравните качество полученных моделей.
5. Произведите для каждой модели подбор значений одного гиперпараметра. В зависимости от используемой библиотеки можно применять функцию `GridSearchCV`, использовать перебор параметров в цикле, или использовать другие методы.
6. Повторите пункт 4 для найденных оптимальных значений гиперпараметров. Сравните качество полученных моделей с качеством моделей, полученных в пункте 4.

Lab6

June 3, 2019

```
In [1]: !pip install seaborn
        !pip install lightgbm
```

```
Requirement already satisfied: seaborn in /srv/conda/lib/python3.6/site-packages (0.9.0)
Requirement already satisfied: pandas>=0.15.2 in /srv/conda/lib/python3.6/site-packages (from s
Requirement already satisfied: matplotlib>=1.4.3 in /srv/conda/lib/python3.6/site-packages (fr
Requirement already satisfied: numpy>=1.9.3 in /srv/conda/lib/python3.6/site-packages (from se
Requirement already satisfied: scipy>=0.14.0 in /srv/conda/lib/python3.6/site-packages (from s
Requirement already satisfied: python-dateutil>=2.5.0 in /srv/conda/lib/python3.6/site-packages
Requirement already satisfied: pytz>=2011k in /srv/conda/lib/python3.6/site-packages (from pan
Requirement already satisfied: cycycler>=0.10 in /srv/conda/lib/python3.6/site-packages (from ma
Requirement already satisfied: kiwisolver>=1.0.1 in /srv/conda/lib/python3.6/site-packages (fr
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /srv/conda/lib/pyth
Requirement already satisfied: six>=1.5 in /srv/conda/lib/python3.6/site-packages (from python
Requirement already satisfied: setuptools in /srv/conda/lib/python3.6/site-packages (from kiwi
Requirement already satisfied: lightgbm in /srv/conda/lib/python3.6/site-packages (2.2.3)
Requirement already satisfied: numpy in /srv/conda/lib/python3.6/site-packages (from lightgbm)
Requirement already satisfied: scipy in /srv/conda/lib/python3.6/site-packages (from lightgbm)
Requirement already satisfied: scikit-learn in /srv/conda/lib/python3.6/site-packages (from li
```

```
In [2]: import numpy as np
import pandas as pd
from sklearn import datasets
import matplotlib.pyplot as plt
import seaborn as sns
import lightgbm

from sklearn.metrics import accuracy_score, balanced_accuracy_score
from sklearn.metrics import precision_score, recall_score, f1_score, classification_re
from sklearn.metrics import confusion_matrix
from sklearn.metrics import mean_absolute_error, mean_squared_error, mean_squared_log_
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.neighbors import KNeighborsRegressor, KNeighborsClassifier
from sklearn.model_selection import learning_curve, validation_curve
```

```

%matplotlib inline

In [3]: df = pd.read_csv('heart.csv', sep=",")
import warnings
warnings.filterwarnings("ignore")

```

1 Data Set Information:

This database contains 76 attributes, but all published experiments refer to using a subset of 14 of them. In particular, the Cleveland database is the only one that has been used by ML researchers to this date. The "goal" field refers to the presence of heart disease in the patient. It is integer valued from 0 (no presence) to 4.

1.1 Attribute Information:

1. age
2. sex
3. chest pain type (4 values)
4. resting blood pressure
5. serum cholestoral in mg/dl
6. fasting blood sugar > 120 mg/dl
7. resting electrocardiographic results (values 0,1,2)
8. maximum heart rate achieved
9. exercise induced angina
10. oldpeak = ST depression induced by exercise relative to rest
11. the slope of the peak exercise ST segment
12. number of major vessels (0-3) colored by flourosopy
13. thal: 3 = normal; 6 = fixed defect; 7 = reversable defect

```
In [4]: df.shape
```

```
Out[4]: (303, 14)
```

```
In [5]: df.head()
```

```
Out[5]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
In [6]: df.dtypes
```

```
Out[6]: age          int64
sex            int64
cp             int64
trestbps       int64
chol           int64
fbs            int64
restecg        int64
thalach         int64
exang          int64
oldpeak        float64
slope          int64
ca             int64
thal           int64
target         int64
dtype: object
```

```
, :
```

```
In [7]: df.isnull().sum()
```

```
Out[7]: age          0
sex            0
cp             0
trestbps       0
chol           0
fbs            0
restecg        0
thalach         0
exang          0
oldpeak        0
slope          0
ca             0
thal           0
target         0
dtype: int64
```

```
.
:
```

```
In [8]: df['target'].unique()
```

```
Out[8]: array([1, 0])
```

```
-,
```

```
In [9]: #sns.pairplot(df, hue= "target")
```

```
In [10]: plt.figure(figsize=(15, 10))
sns.heatmap(df.corr(method='spearman'),annot=True, fmt='.3f')
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7c633efd30>
```

```
Out[10]:
```



```
In [11]: #lgbm_regressor = lightgbm.LGBMRegressor().fit(df.loc[:, df.columns != 'target'], df[
#list_of_importances = list(zip(df.loc[:, df.columns != 'target'].columns.tolist(),
#lgbm_regressor.feature_importances_))

#sorted(list_of_importances, key= lambda x: x[1], reverse= True)
```

```
In [12]: ##important_features = [x[0] for x in sorted(list_of_importances, key= lambda x: x[1]
```

1.2 .

```
In [13]: X_train, X_test, y_train, y_test = train_test_split(df.loc[:, df.columns != 'target'],
df['target'],
test_size= 0.2,
random_state= 42)
```

```
In [45]: from sklearn.ensemble import RandomForestClassifier
```

2

```
In [46]: clas = RandomForestClassifier().fit(X_train, y_train)

In [47]: print(accuracy_score(clas.predict(X_test), y_test))
         print(f1_score(clas.predict(X_test), y_test, average='macro'))
         print(precision_score(clas.predict(X_test), y_test, average='macro'))

0.819672131147541
0.8196721311475409
0.8216594827586207
```

```
In [49]: from sklearn.ensemble import GradientBoostingClassifier
```

3

```
In [50]: clas=GradientBoostingClassifier(random_state=1)
         clas.fit(X_train, y_train)

Out[50]: GradientBoostingClassifier(criterion='friedman_mse', init=None,
                                     learning_rate=0.1, loss='deviance', max_depth=3,
                                     max_features=None, max_leaf_nodes=None,
                                     min_impurity_decrease=0.0, min_impurity_split=None,
                                     min_samples_leaf=1, min_samples_split=2,
                                     min_weight_fraction_leaf=0.0, n_estimators=100,
                                     n_iter_no_change=None, presort='auto', random_state=1,
                                     subsample=1.0, tol=0.0001, validation_fraction=0.1,
                                     verbose=0, warm_start=False)

In [51]: print(accuracy_score(clas.predict(X_test), y_test))
         print(f1_score(clas.predict(X_test), y_test, average='macro'))
         print(precision_score(clas.predict(X_test), y_test, average='macro'))

0.7704918032786885
0.7704301075268818
0.771551724137931
```

```
In [23]: from sklearn.model_selection import GridSearchCV
```

4 - n_estimators

```
In [52]: clas = RandomForestClassifier()
         param = {'n_estimators':range(10,300,10)}
         GV = GridSearchCV(clas, param, cv=3)
         GV.fit(X_train, y_train)
```

```

Out [52]: GridSearchCV(cv=3, error_score='raise-deprecating',
                      estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criterion=
                        max_depth=None, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators='warn', n_jobs=None,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False),
                      fit_params=None, iid='warn', n_jobs=None,
                      param_grid={'n_estimators': range(10, 300, 10)},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                      scoring=None, verbose=0)

```

```

In [53]: GV.best_estimator_

```

```

Out [53]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=160, n_jobs=None,
                                oob_score=False, random_state=None, verbose=0,
                                warm_start=False)

```

```

In [54]: print(accuracy_score(GV.predict(X_test), y_test))

```

```

0.8524590163934426

```

```

In [62]: clas = GradientBoostingClassifier()
          param = {'n_estimators': range(10,500,10)}
          GV = GridSearchCV(clas, param, cv=3)
          GV.fit(X_train, y_train)

```

```

Out [62]: GridSearchCV(cv=3, error_score='raise-deprecating',
                      estimator=GradientBoostingClassifier(criterion='friedman_mse', init=None,
                        learning_rate=0.1, loss='deviance', max_depth=3,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_sampl... subsample=1.0, tol=0.0001, valida
                        verbose=0, warm_start=False),
                      fit_params=None, iid='warn', n_jobs=None,
                      param_grid={'n_estimators': range(10, 500, 10)},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                      scoring=None, verbose=0)

```

```

In [63]: GV.best_estimator_

```

```

Out [63]: GradientBoostingClassifier(criterion='friedman_mse', init=None,
                                      learning_rate=0.1, loss='deviance', max_depth=3,

```



```
max_features=None, max_leaf_nodes=None,  
min_impurity_decrease=0.0, min_impurity_split=None,  
min_samples_leaf=1, min_samples_split=2,  
min_weight_fraction_leaf=0.0, n_estimators=50,  
n_iter_no_change=None, presort='auto', random_state=None,  
subsample=1.0, tol=0.0001, validation_fraction=0.1,  
verbose=0, warm_start=False)
```

```
In [64]: print(accuracy_score(GV.predict(X_test), y_test))
```

```
0.8360655737704918
```

```
In [0]:
```